

CS479: Machine Learning for 3D Data

3D Representations

LECTURE 2
MINHYUK SUNG

Spring 2025
KAIST

Previously in CS479

Applications: 3D Reconstruction $\rightarrow$ decoding net.




Google Maps Immersive View

<https://cdn.mos.cms.futurecdn.net/P7HseGaXpSTQM2uAfSbh5Y.jp>

g

Previously in CS479

Applications: 3D Generation $\rightarrow$ *decoding Net.*




Wang et al., ProlificDreamer: High-Fidelity and Diverse Text-to-3D Generation with Variational Score Distillation, arXiv 2023.



Roblox

3D Encoder


A neural network taking 3D data as **input**.



<https://towardsdatascience.com/applied-deep-learning-part-3-autoencoders-1c083af4d798>

3D Decoder

A neural network generating 3D data as **output**.



<https://towardsdatascience.com/applied-deep-learning-part-3-autoencoders-1c083af4d798>

Course Road Map

Representations


- Point clouds
- Implicit representation
- Multi-view images to 3D
- Hybrid representations
- Meshes
- ~~CAD~~
- Representation Conversion

Applications


- 3D perception
- Reconstruction
- Manipulation
- ~~Generation~~

3D Representations

Texts




Images




```
[[105 112 108 111 104 99 106 99 96 103 112 119 104 97 93 87]
 [ 91 98 102 106 104 79 98 103 99 105 123 136 110 105 94 85]
 [ 76 85 90 105 128 105 87 96 95 99 115 112 106 103 99 85]
 [ 99 81 81 93 120 131 127 100 95 98 102 99 96 93 101 94]
 [106 91 61 64 69 91 88 85 101 107 109 98 75 84 96 95]
 [114 108 85 55 55 69 64 54 64 87 112 129 98 74 84 91]
 [133 137 147 103 65 81 80 65 52 54 74 84 102 93 85 82]
 [128 137 144 140 109 95 86 70 62 65 63 63 60 73 86 101]
 [125 133 148 137 119 121 117 94 65 79 80 65 54 64 72 98]
 [127 125 131 147 133 127 126 131 111 96 89 75 61 64 72 84]
 [115 114 109 123 150 148 131 118 113 109 100 92 74 65 72 78]
 [ 89 93 90 97 108 147 131 118 113 114 113 109 106 95 77 80]
 [ 63 77 86 81 77 79 102 123 117 115 117 125 125 130 115 87]
 [ 62 65 82 89 78 71 80 101 124 126 119 101 107 114 131 119]
 [ 63 65 75 88 89 71 62 81 120 138 135 105 81 98 110 118]
 [ 87 65 71 87 106 95 69 45 76 130 126 107 92 94 105 112]
 [118 97 82 86 117 123 116 66 41 51 95 93 89 95 102 107]
 [164 146 112 80 82 120 124 104 76 48 45 66 88 101 102 109]
 [157 170 157 120 93 86 114 132 112 97 69 55 70 82 99 94]
 [130 128 134 161 139 100 109 118 121 134 114 87 65 53 69 86]
 [128 112 96 117 150 144 120 115 104 107 102 93 87 81 72 79]
 [123 107 96 86 83 112 153 149 122 109 104 75 80 107 112 99]
 [122 121 102 80 82 86 94 117 145 148 153 102 58 78 92 107]
 [122 164 148 103 71 56 78 83 93 103 119 139 102 61 69 84]]
```


Convolutional Neural Network



Transformers




Transformer Encoder



3D Grid (Voxels)


Voxel : volume(부피) + Element
; 3차원 공간의 pixel.

3D Grid



<https://towardsdatascience.com/a-comprehensive-introduction-to-different-types-of-convolutions-in-deep-learning-669281e58215>

Medical Imaging




Mathotaarachchi et al., VoxelStats, 2016.

We're interested in 2D surfaces in 3D space.


9 3d info from 3d scan. or 3d model. → issue: volume info에 관심 X. surface info에 관심 O.

9 we don't know and care about inside info.

→ scene object들 이해. (3d space interior empty!)



The Digital Michelangelo Project



Liu et al., A Local/Global Approach to Mesh Parameterization, SGP 2008.

How we encode all 3d info into the voxels

→ one way is encoding occupancy info like which voxel is inside or which voxel is outside.

Encoding Surface Info

→ checking what are the voxels where the surface is basically passing.

Encode voxel having the surface is 1.

other voxels are 0. (empty voxel 이니까).

2d vs 3d.

2d: 해상도 2배 up → 전체 픽셀 수 $2 \times 2 \times 2 = 4$ 배 up.

3d: // → 전체 Voxel 수 $2 \times 2 \times 2 = 8$ 배 up.

* 전체 부피 계산과 표면 구상 부피 계산 늘리는 속도가 다름.

공간 부피에 비례.

$$O(N^3)$$

물체 겉넓이에 비례.

$$O(N^2)$$

$$\frac{\text{Surface Voxels}}{\text{Total Voxels}} \approx \frac{N^2}{N^3} = \frac{1}{N}.$$

해상도(N)가 커질수록 전체 공간에서 표면 voxel이 차지하는 비율은 $1/N$ 에 비례해서 계속 작아짐.

"고해상도로 간혹 객체 있는 '표면'. 데이터는


3D Convolution

전체 공간의 아주 극소수에만 존재. 99%는 내부/외부 empty space나 태는 쓸모없는 voxel."

- Each voxel contains a binary value (indicating whether the voxel is on the surface or not), and also, most voxels are empty. → whole representation becomes very inefficient and redundant."
- Huge waste of computation.



실제 표면 정보는
가지고 있는
voxel



해상도 높일수록
표면에 있는 voxel 수가
sparse 해짐.



(sparse representation)



3D CNNs

(-) Takes a huge amount of memory and time in training.


↳ 3d volume은 기본적으로 very sparse information을 가지고 있음.




3D CNNs

Bigger voxels in the empty space.

Architectures using adaptive data structure



Wang et al., O-CNN, SIGGRAPH 2017.



Riegler et al., O-CNN, CVPR 2017.





Image from Nvidia



<https://miro.medium.com/>

a no surface info
→ bigger voxel.


a increasing resolution only in
some kind of local region
where we know that the surface is basically located.

convolution 연산에서
adaptive resolution
을 사용해서
깊이 더 내려가면
해상도 lower

SparseConvNet [Grahamet al., 2018]

- Compute convolutions only in the active areas.
- Still takes lots of memory and time in training.

surface가 located 된 region만.




Multi-View Images

2015부터는 모든 3d data를 이미지 set으로 처리하는 방식 등장~.

↳ Multi-View Image representation 시도한 것.

Multi-View CNN [Su et al., 2015]



3D Object Classification

	input	#views	accuracy avg. class	accuracy overall
SPH [12]	mesh	-	68.2	-
3DShapeNets [29]	volume	1	77.3	84.7
VoxNet [18]	volume	12	83.0	85.9
Subvolume [19]	volume	20	86.0	89.2
LFD [29]	image	10	75.5	-
MVCNN [24]	image	80	90.1	-

3d conv

2d conv.

multi-view.


Dataset: ModelNet40

Metric: 40-class classification accuracy (%)

3d data 처리를 위한 special

Net 이 필요함!!

3D Segmentation with 2D Neural Networks




Multi-View Images

- (+) Especially good for processing **appearance** information like color, texture, and material.
- (−) Requires lots of images for high accuracy and thus takes lots of memory and time.
- (−) May not be able to capture geometric details.

Point Cloud



Point Cloud

- The simplest representation: **only points**, no connectivity.
- Collection of (x,y,z) coordinates, possibly with normal.



Point Cloud

- Nearly all 3D scanning devices produce point clouds.




<https://techbullion.com/wp-content/uploads/2021/12/Mobile-LiDAR-Scanner.jpg>

Point Cloud


- Nearly all 3D scanning devices produce point clouds.
- Sometimes, easier to handle.

Fracturing Solids



Meshless Animation of Fracturing Solids
Pauly et al., SIGGRAPH '05


Fluids




Adaptively sampled particle fluids,
Adams et al. SIGGRAPH '07

Neural Networks for Point Clouds

- (+) Fast, easy to implement (relatively).
- (+) The most popular architectures.




Qi et al., PointNet, CVPR 2017.



Qi et al., PointNet++, NeurIPS 2017.


Point Cloud

- (—) No surface/topology information;
needs to be converted to the other representation for downstream applications.



Point Cloud

- (—) No surface/topology information; needs to be converted to the other representation for downstream applications.
- (—) Weak approximation power; requires many points for the details.



Li et al., PU-GAN, ICCV 2019.

Polygon Mesh

Polygon Mesh


- The **most popular** representation for shapes in graphics.
- A compact form representing **surfaces**.
- A **graph-like** structure but not the same.



<https://www.3dcadbrowser.com/3d-model/people-collection-low-poly-62652>

Polygon Mesh

- A polygon mesh is a collection of **vertices**, **edges** and **faces** that defines the shape of a polyhedral object.
- A triangle mesh is a special case when all the faces are **triangles**.




$$M = \langle V, E, F \rangle$$

vertices


edges

faces


Online Repositories of 3D Meshes



ShapeNet



3D Warehouse



Yobi 3D




SceneNN




Redwood Dataset



ScanNet




KITTI 3D



Dynamic MPI-FAUST


Polygon Mesh

- (+) Good for many applications:
 - Rendering




Polygon Mesh

- (+) Good for many applications:
 - Rendering
 - Texturing




<https://dreamfarmstudios.com/blog/getting-to-know-3d-texturing-in-animation-production/>



https://commons.wikimedia.org/wiki/File:Displacement_Mapping.jpg

Polygon Mesh

- (+) Good for many applications:
 - Rendering
 - Texturing
 - Deformation / Manipulation




Polygon Mesh

- (+) Good for many applications:
 - Rendering
 - Texturing
 - Deformation / Manipulation
 - Simulation




<https://3dmodelsworld.com/maya-bullet-physics-simulation-tutorial-wrecking-ball-animation-active-and-passive-rigid-body/>




<https://www.youtube.com/watch?v=HKL8mQ01iuU>

Polygon Mesh

- (+) Good for many applications:
 - Rendering
 - Texturing
 - Deformation / Manipulation
 - Simulation
- (−) Irregular structure



<https://towardsdatascience.com/convolution-vs-correlation-af868b6b4fb5>



Alliez et al., Recent Advances in Remeshing of Surfaces.

Pooling in Neural Network


Aggregating information while progressively reducing the resolution of the data.

12	20	30	0
8	12	2	0
34	70	37	4
112	100	25	12

2×2 Max-Pool

20	30
112	37

<https://paperswithcode.com/method/max-pooling>



Gu et al., Blind Channel Identification Aided Generalized Automatic Modulation Recognition Based on Deep Learning, 2019.


Pooling Operation for Polygon Mesh

How can the resolution of a polygon mesh be decreased?

Through the process of iterative **edge contraction**.




https://doc.cgal.org/latest/Surface_mesh_simplification/index.html



Yao et al., Quadratic Error Metric Mesh Simplification
Algorithm Based on Discrete Curvature, 2015


MeshCNN [Hanoka et al., 2019]

Pool adjacent edge information via the **edge contraction**.




Neural Networks for Meshes


- Requires parameterization or specialized operations.
- Hard to implement. Verified only with a few use cases.




Masci et al., GCNN, ICCV 2015.



Hanocka et al., MeshCNN, SIGGRAPH 2019.



Mitchel et al., ICCV 2021.



Milano et al., NeurIPS 2020.


Polygon Mesh

- (+) Good for many applications:
 - Rendering
 - Texturing
 - Deformation / Manipulation
 - Simulation
- (−) Irregular structure
- (−) Difficult to create a valid mesh

Valid Meshes

E.g. 2-manifoldness

Each local region should be homeomorphic (mappable) to a 2D flat plane.




<https://www.shapeways.com/blog/archives/29453-tutorial-tuesday-5-quick-fixes-with-meshlab.html>

Valid Meshes

- Watertightness
- Topology
- Normal orientation consistency


Etc.



<https://courses.cs.duke.edu/fall06/cps296.1/Lectures/sec-II-1.pdf>


Mesh Generation

E.g., an autoregressive model generating vertices and faces sequentially.



Mesh Generation

E.g., an autoregressive model generating vertices and faces
sequentially.




Point-Cloud Conditioned Mesh Generation Progress


CAD Representations

CAD Representations


- More compact and structural representations.
- NURBS, CSG, B-Rep, Extrusions, Revolve, etc.





Autodesk



CSG




Karl D.D. Willis




steemit

CAD Representations


- Very few neural networks processing and generating them.



Lambourne et al., BRepNet, CVPR 2021.




Jayaraman et al., UV-Net, CVPR 2021.




Ren et al., CSG-Stump, ICCV 2021.

CAD Representations

- Very few neural networks processing and generating them.
- The architecture are very complex.



3D Representations




	Voxels	Multi-View	Point Cloud	Mesh
Pros	• Okay for processing if we use sparse convolution	• Good for processing • Easy to implement	• Good for both processing and generation • Efficient • Easy to implement	• Compact • Good for many applications
Cons	• Hard to implement • Inefficient	• Inefficient	• Require lots of points to capture details	• Bad for both processing and generation


Implicit Representation

Implicit Representation

- A **function** that takes coordinates as input and returns occupancy or signed distance.
- Representation for **output (generation)** not input (processing).




Park et al., DeepSDF, CVPR 2020.



Mescheder et al., Occupancy Networks, CVPR 2020.


3D Representations



	Voxels	Multi-View	Point Cloud	Mesh	Implicit
Pros	• Okay for processing if we use sparse convolution	• Good for processing • Easy to implement	• Good for both processing and generation • Efficient • Easy to implement	• Compact • Good for many applications	• Good for generation
Cons	• Hard to implement • Inefficient	• Inefficient	• Require lots of points to capture details	• Bad for both processing and generation	• Cannot be used for processing • Need conversion for applications

Conversion Across Representations

Conversion Across Representations



Course Road Map

Representations

- **Point clouds**
- Implicit representation
- Multi-view images to 3D
- Hybrid representations
- Meshes
- ~~CAD~~
- Representation Conversion

Applications

- **3D perception (Encoding)**
- **Reconstruction (Decoding)**
- Manipulation
- ~~Generation~~